

# Linear\_Regression

December 7, 2025

## Simple Linear Regression

- In statistics, simple linear regression is a linear regression model with a single explanatory variable.
- In simple linear regression, we predict scores on one variable based on results on another.
- The criteria variable Y is the variable we are predicting. Predictor variable X is the variable using which we are making our predictions.
- The prediction approach is known as simple regression as there is only one predictor variable.
- As a result, a linear function that predicts the values of the dependent variable as a function of the independent variable is discovered for two-dimensional sample points with one independent variable and one dependent variable.

**Equation :**  $y = \beta_0 + \beta_1 x$

- This is the simple linear regression equation where  $\beta_0$  is the constant and  $\beta_1$  is the slope and describes the relationship between  $x$  (independent variable) and  $y$  (dependent variable)
- $y_i = \beta_0 + \beta_1 x_i$ ,  $\beta_0$  ( $y$ -intercept) and  $\beta_1$  (slope) are the coefficients whose values represent the accuracy of predicted values with the actual values.

**Implementation** In this example, we will use the `salary_Data` concerning the experience of employees. In this dataset, we have two columns `YearsExperience` and `Salary`

```
read.csv('Salary_Data.csv')
```

```
[1]: # Import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
[2]: # Get dataset
df_sal = pd.read_csv('Salary_Data.csv')
df_sal.head()
```

```
[2]:   YearsExperience    Salary
0            1.1  39343.0
1            1.3  46205.0
2            1.5  37731.0
```

```
3           2.0  43525.0  
4           2.2  39891.0
```

```
[3]: # Describe data  
df_sal.describe()
```

```
[3]:      YearsExperience      Salary  
count      30.000000      30.000000  
mean       5.313333    76003.000000  
std        2.837888   27414.429785  
min        1.100000   37731.000000  
25%        3.200000   56720.750000  
50%        4.700000   65237.000000  
75%        7.700000  100544.750000  
max       10.500000  122391.000000
```

- Here, we can see Salary ranges from 37731 to 122391 and a median of 65237

### Relationship between Salary and Experience

```
[4]: plt.scatter(df_sal['YearsExperience'], df_sal['Salary'], color ='lightcoral')  
plt.title('Salary vs Experience')  
plt.xlabel('Years of Experience')  
plt.ylabel('Salary')  
plt.box(True)  
plt.show()
```



- It is clearly visible now, our data varies linearly.
- That means, that an individual receives more **Salary** as they gain **Experience**.

### Split the dataset into dependent/independent variables

```
[5]: # Splitting variables
x = df_sal.iloc[:, :1]
# independent
y = df_sal.iloc[:, 1:]
# dependent
```

**Split data into Train/Test sets** Further, split your data into training (80%) and test (20%) sets using `train_test_split`

```
[6]: # Splitting dataset into test/train
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2,
random_state = 0)
```

### Train the regression model

```
[7]: # Regressor model
regressor = LinearRegression()
```

```
regressor.fit(x_train, y_train)
```

[7]: LinearRegression()

**Predict the result** Here comes the interesting part, when we are all set and ready to predict any value of y (Salary) dependent on x (Experience) with the trained model using regressor.predict

[8]: # Prediction result  
y\_pred\_test = regressor.predict(x\_test) # predicted value of y\_test  
y\_pred\_train = regressor.predict(x\_train) # predicted value of y\_train

### Plot training set data vs predictions

- First we plot the result of training sets (x\_train, y\_train) with x\_train and predicted value of y\_train (regressor.predict(x\_train))

[9]: # Prediction on training set  
plt.scatter(x\_train, y\_train, color = 'lightcoral')  
plt.plot(x\_train, y\_pred\_train, color = 'firebrick')  
plt.title('Salary vs Experience (Training Set)')  
plt.xlabel('Years of Experience')  
plt.ylabel('Salary')  
plt.legend(['x\_train/y\_train', 'x\_train/Pred(y\_test)'], loc='best',  
 facecolor='white')  
plt.box(True)  
plt.show()



**Plot test set data vs predictions** Secondly, we plot the result of test sets (`X_test`, `y_test`) with `X_train` and predicted value of `y_train` (`regressor.predict(X_train)`)

```
[10]: # Prediction on test set
plt.scatter(x_test, y_test, color = 'lightcoral')
plt.plot(x_train, y_pred_train, color = 'firebrick')
plt.title('Salary vs Experience (Test Set)')
plt.xlabel('Years of Experience')
plt.ylabel('Salary')
plt.legend(['x_train/y_train', 'x_train/Pred(y_test)'], loc='best',
          facecolor='white')
plt.box(True)
plt.show()
```



```
[11]: # Regressor coefficients and intercept
print(f'Coefficient beta_1: {regressor.coef_[1]}')
print(f'Intercept beta_0: {regressor.intercept_[0]}' )
```

```
Coefficient beta_1: [[9312.57512673]]
Intercept beta_0: [26780.09915063]
```

```
[ ]:
```